



DEEP
LEARNING
INSTITUTE



PRAIRIE VIEW
A&M UNIVERSITY

DLI Accelerated Data Science Teaching Kit

Lecture 11.3 - Spark SQL and Other Spark Libraries



The Accelerated Data Science Teaching Kit is licensed by NVIDIA, Georgia Institute of Technology, and Prairie View A&M University under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).

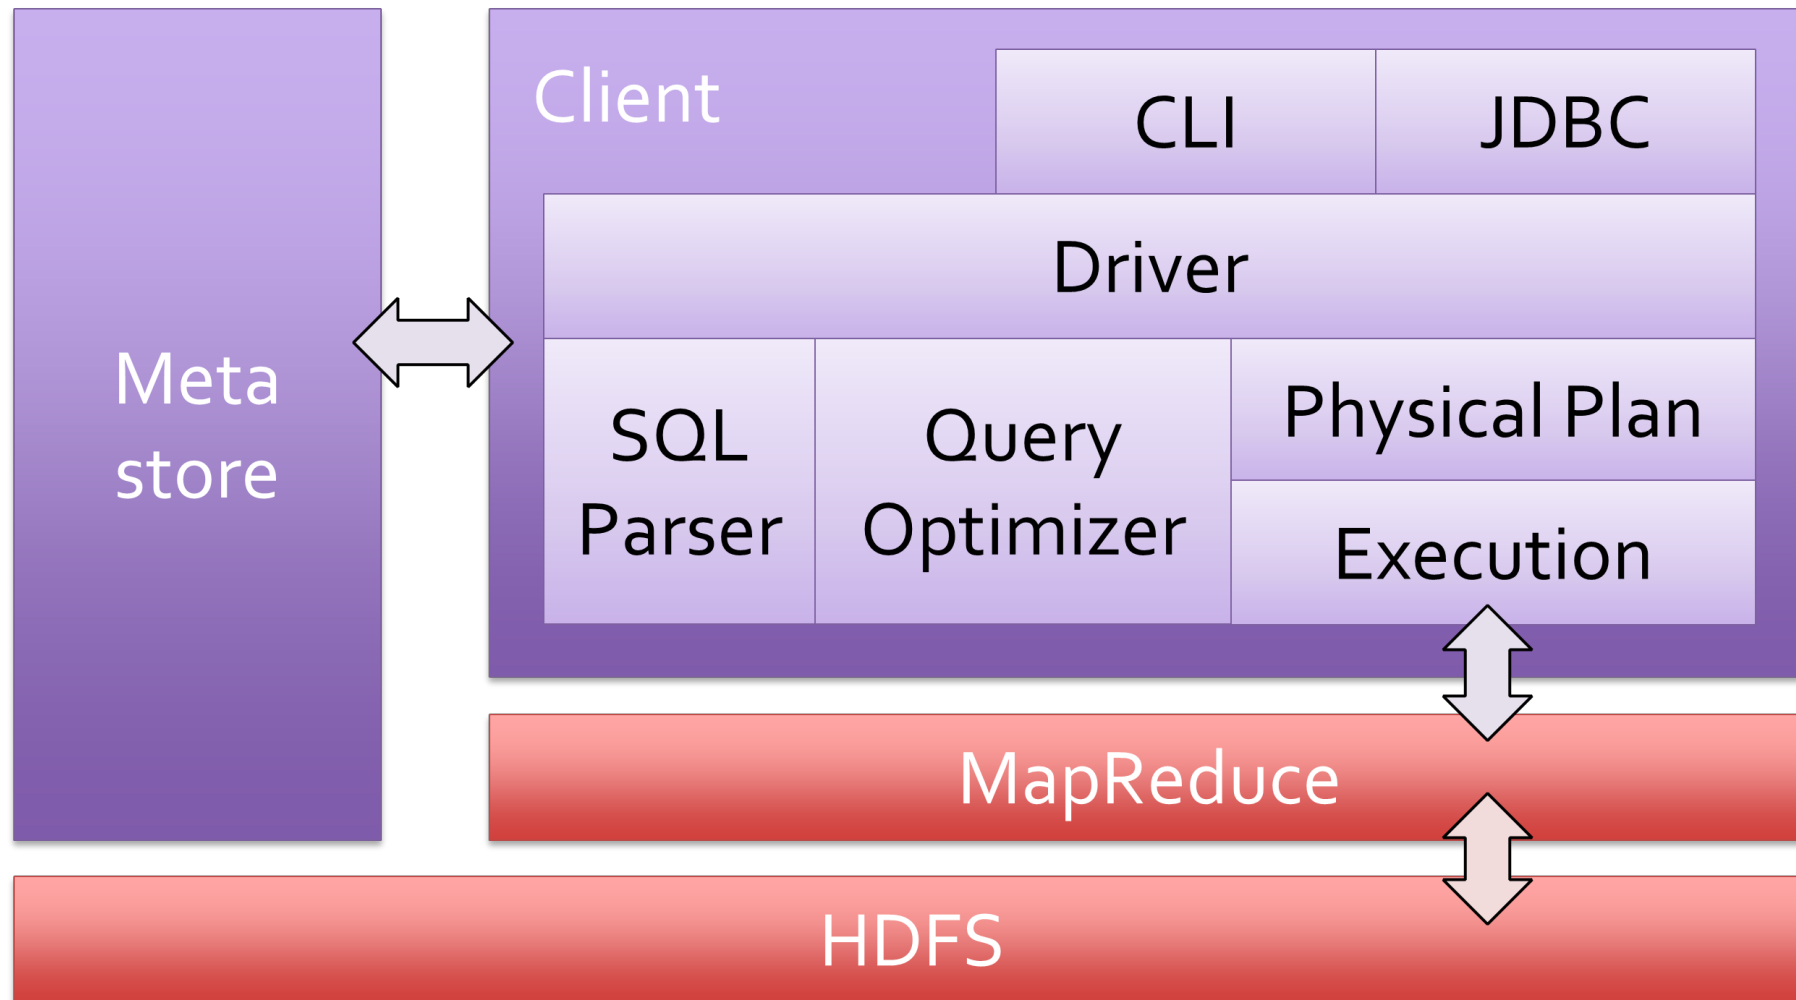
Motivation

Hive is great, but Hadoop's disk-based engine can make even the smallest queries take minutes

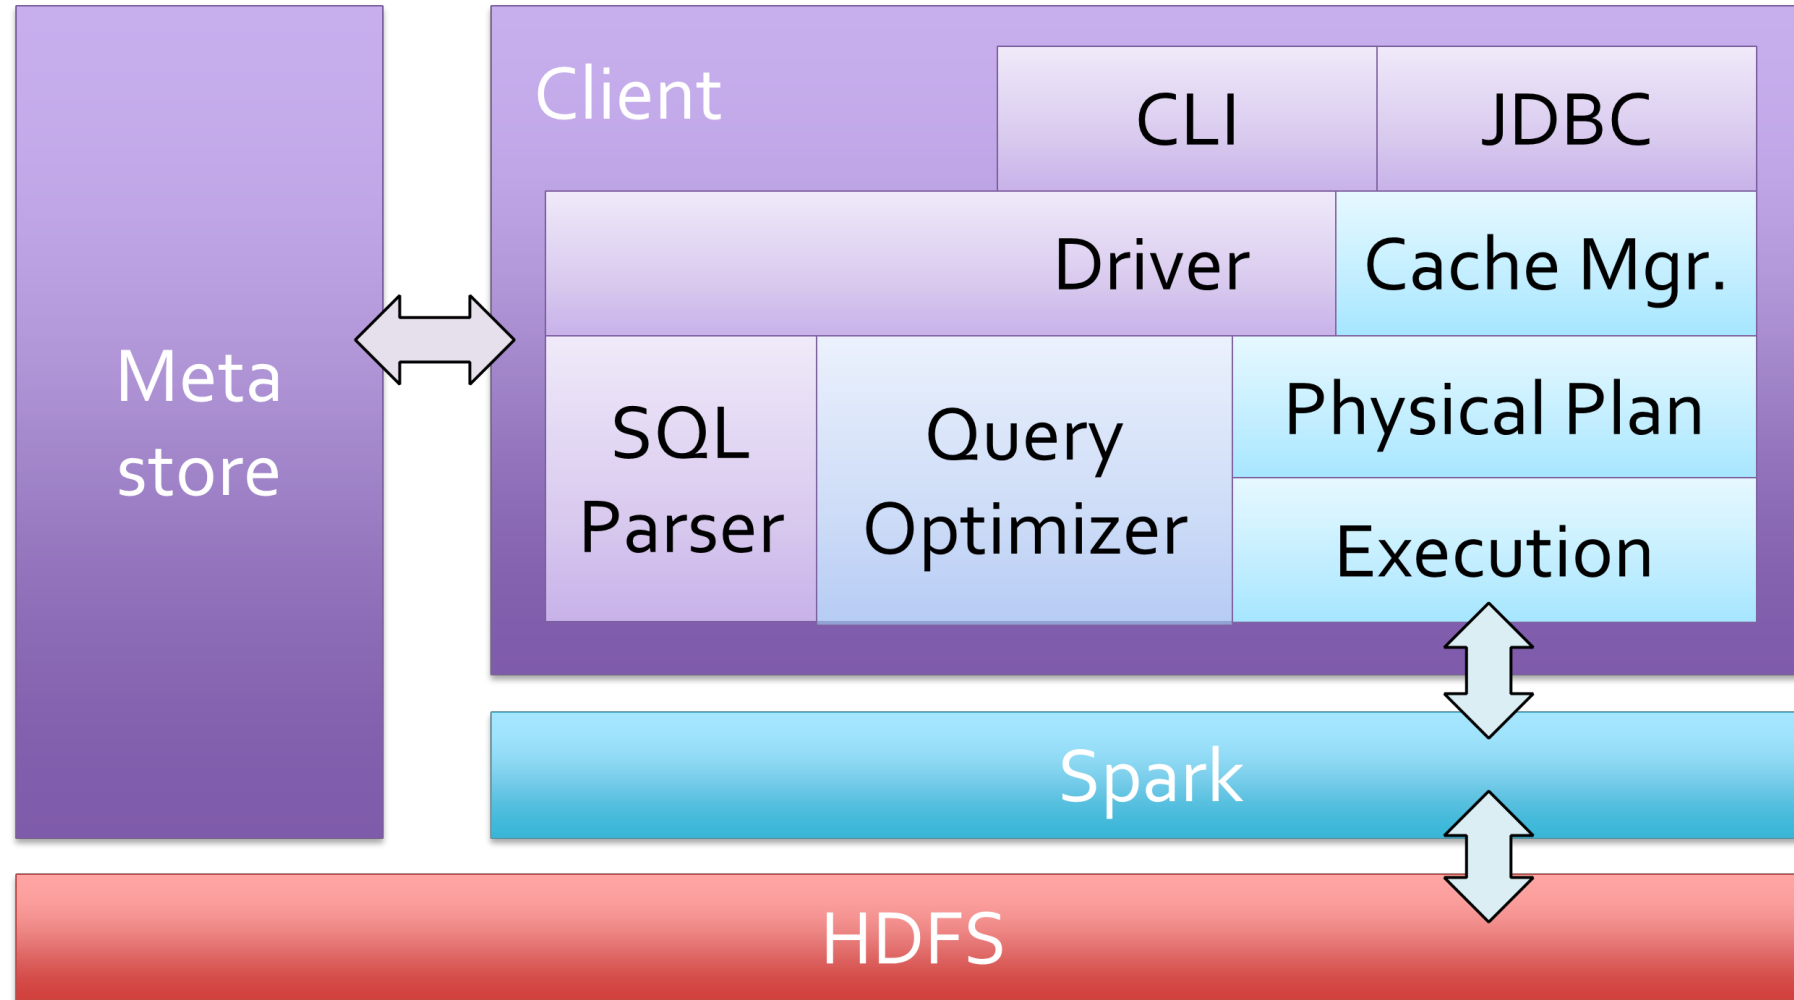
Can we extend Hive to run on Spark?

Yes! Spark SQL = Hive on Spark

Hive Architecture




Spark SQL Architecture



[Engle et al, SIGMOD 2012]


Benchmark Query 1

SELECT * FROM grep WHERE field LIKE '%XYZ%';



Benchmark Query 2

```
SELECT sourceIP, AVG(pageRank), SUM(adRevenue) AS earnings
FROM rankings AS R, userVisits AS V ON R.pageURL = V.destURL
WHERE V.visitDate BETWEEN '1999-01-01' AND '2000-01-01'
GROUP BY V.sourceIP
ORDER BY earnings DESC
LIMIT 1;
```



Spark Streaming

Recall that Spark's model was motivated by two emerging uses (interactive and multi-stage apps)

Another emerging use case that needs fast data sharing is **stream processing**

- » Track and update state in memory as events arrive
- » Large-scale reporting, click analysis, spam filtering, etc.


Spark Streaming

Extends Spark to perform streaming computations

Runs as a series of small (~1 s) batch jobs, keeping state in memory as fault-tolerant RDDs

Intermix seamlessly with batch and ad-hoc queries

```
tweetStream  
  .flatMap(_.toLowerCase.split)  
  .map(word => (word, 1))  
  .reduceByWindow(5, _ + _)
```



[Zaharia et al, HotCloud 2012]

Spark Streaming

Extends Spark to perform streaming computations

Runs as a series of small (~1 s) batch jobs,
keeping state in memory as fault-tolerant RDDs

Intermix seamlessly with batch and ad-hoc queries



Result: can process **42 million** records/second
(4 GB/s) on 100 nodes at **sub-second** latency

[Zaharia et al, HotCloud 2012]

MLlib

- Basic statistics
 - summary statistics
 - correlations
 - stratified sampling
 - hypothesis testing
 - streaming significance testing
 - random data generation
- Classification and regression
 - linear models (SVMs, logistic regression, linear regression)
 - naive Bayes
 - decision trees
 - ensembles of trees (Random Forests and Gradient-Boosted Trees)
 - isotonic regression
- Collaborative filtering
 - alternating least squares (ALS)
- Clustering
 - k-means
 - Gaussian mixture
 - power iteration clustering (PIC)
 - latent Dirichlet allocation (LDA)
 - bisecting k-means
 - streaming k-means
- Dimensionality reduction
 - singular value decomposition (SVD)
 - principal component analysis (PCA)
- Feature extraction and transformation
- Frequent pattern mining
 - FP-growth
 - association rules
 - PrefixSpan
- Evaluation metrics
- PMML model export
- Optimization (developer)
 - stochastic gradient descent
 - limited-memory BFGS (L-BFGS)

<https://spark.apache.org/docs/latest/ml-guide.html>

GraphX


Parallel graph processing

Extends RDD -> Resilient Distributed Property Graph


- » Directed multigraph with properties attached to each vertex and edge

Limited algorithms

- » PageRank
- » Connected Components
- » Triangle Counts



Behavior with Not Enough RAM





DEEP
LEARNING
INSTITUTE



DLI Accelerated Data Science Teaching Kit

Thank You

We thank Dr. Matei Zaharia for sharing teaching materials for Spark.