



DEEP
LEARNING
INSTITUTE



DLI Accelerated Data Science Teaching Kit

Lecture 12.2 - How HBase Scales Up Storage



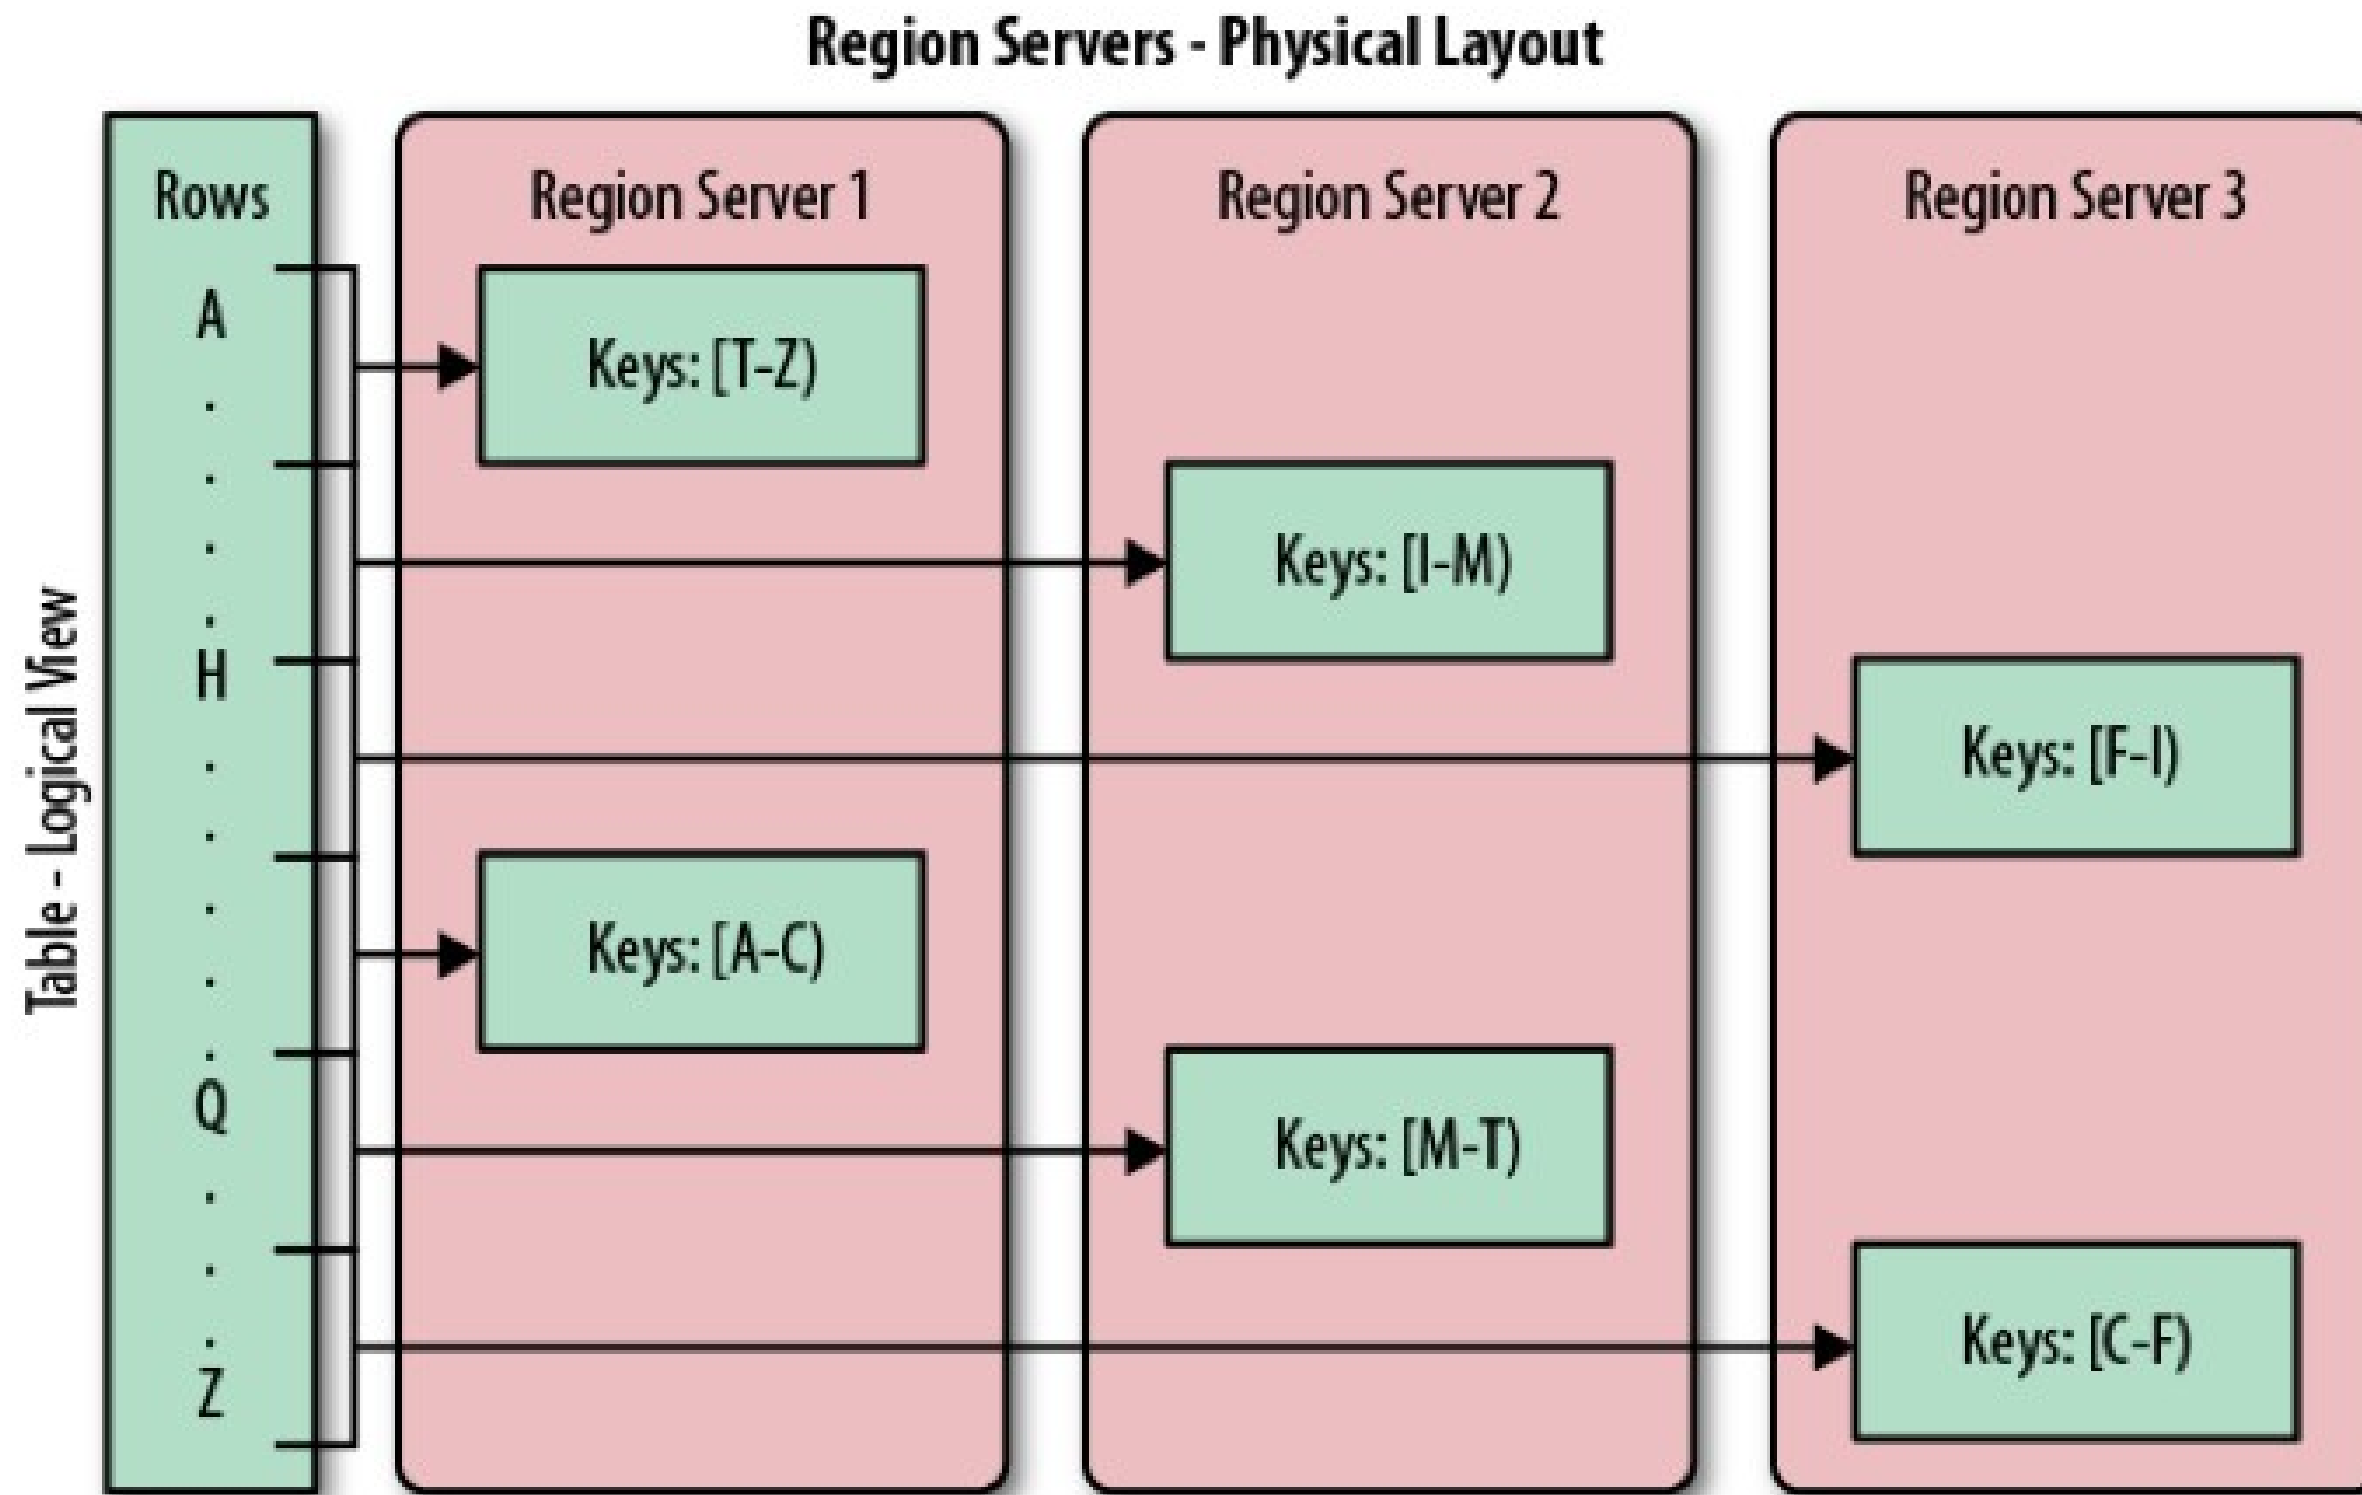
The Accelerated Data Science Teaching Kit is licensed by NVIDIA, Georgia Institute of Technology, and Prairie View A&M University under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).

How HBase scales up storage?

Automatically divide contiguous ranges of rows into regions

Start with one region, split into two when getting too large, and so on.

How HBase scales up storage?



Excellent Summary: <http://blog.cloudera.com/blog/2013/04/how-scaling-really-works-in-apache-hbase/>

Concise way to describe all these?

HBase data model

- Sparse, distributed, persistent, multidimensional map
- Indexed by **row** key + **column** key + **timestamp**

(Table, RowKey, Family, Column, Timestamp) → Value



DEEP
LEARNING
INSTITUTE



PRAIRIE VIEW
A&M UNIVERSITY

DLI Accelerated Data Science Teaching Kit

Thank You